

The IBM logo is displayed in the top left corner of the slide. The background of the slide features a blue-toned image of a computer keyboard with binary code (0s and 1s) overlaid on it.

IBM

Artificial Intelligence Python

Training Module

What is CNN?

Convolutional Neural Networks (CNNs) are a class of deep learning algorithms that are particularly well-suited for analyzing visual data, such as images and videos.

They are inspired by the way the human visual system processes data and images.

CNNs are highly effective in various computer vision tasks because they automatically detect and learn spatial hierarchies in images, such as edges, textures, shapes, and objects.

CNNs are made up of multiple layers, each performing specific tasks, and they are highly efficient for tasks like image classification, object detection, and segmentation.

Basic Components of CNN

1. Convolutional Layer:

- The core building block of a CNN, it applies convolutional filters (kernels) to the input image or feature map to extract patterns such as edges or textures.
- It performs a sliding window operation, where filters slide over the input, calculating dot products between the filter and small regions of the image.

2. Activation Function (ReLU):

- After the convolution operation, an activation function like ReLU (Rectified Linear Unit) is applied to introduce non-linearity. It helps the model to learn complex patterns by enabling it to model non-linear relationships in the data.

3. Pooling Layer:

- Pooling (commonly max pooling) reduces the spatial dimensions (height and width) of the input, which helps reduce the computational complexity and prevents overfitting. It works by selecting the maximum or average value from a small region (e.g., a 2x2 grid) of the image.

4. Fully Connected Layer (Dense Layer):

- After multiple convolutional and pooling layers, the feature maps are flattened into a 1D vector and passed through fully connected layers. This step is similar to the layers in traditional feedforward neural networks and is used for final classification or regression tasks.

5. Softmax / Sigmoid (for output):

- For classification tasks, CNNs typically use a softmax function for multi-class classification or sigmoid for binary classification as the output layer.

Key Uses of CNNs

1. Image Classification:

- CNNs can classify images into predefined categories, such as identifying whether an image contains a cat or a dog. Common datasets include CIFAR-10, MNIST, and ImageNet.

2. Object Detection:

- CNNs are widely used in detecting objects within images, such as locating cars, pedestrians, or animals. Popular architectures like YOLO (You Only Look Once) and Faster R-CNN are designed for real-time object detection.

3. Image Segmentation:

- CNNs can be used for pixel-level image classification, i.e., segmenting an image into regions of interest. For example, U-Net is a CNN architecture used for medical image segmentation.

4. Face Recognition:

- CNNs have been successfully used to recognize faces in images or videos, for example, in security systems or social media platforms.

5. Video Analysis:

- CNNs can also be extended to video, where they are used for tasks such as action recognition or tracking moving objects.

6. Self-Driving Cars:

- CNNs help autonomous vehicles recognize pedestrians, other vehicles, traffic signs, and road conditions.

Implementing CNN using Python (Keras with TensorFlow)

Below is an example of how to implement a simple CNN for image classification using Python and Keras (with TensorFlow as the backend). This example uses the MNIST dataset (handwritten digits) as input.

Steps for Implementation:

1. Install the necessary libraries (if not installed):

```
pip install tensorflow numpy matplotlib
```

2. Import required libraries:

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
```

3. Load and Preprocess the MNIST Dataset:

```
# Load MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Reshape and normalize the data
x_train = x_train.reshape(-1, 28, 28, 1).astype('float32') / 255.0
x_test = x_test.reshape(-1, 28, 28, 1).astype('float32') / 255.0

# One-hot encode the labels
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)
```

4. Build the CNN Model:

```
model = models.Sequential()

# Convolutional Layer 1
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))

# Pooling Layer 1
model.add(layers.MaxPooling2D((2, 2)))

# Convolutional Layer 2
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

# Pooling Layer 2
```

```

model.add(layers.MaxPooling2D((2, 2)))

# Convolutional Layer 3
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

# Flatten and Fully Connected Layer
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))

# Output Layer
model.add(layers.Dense(10, activation='softmax'))

```

5. Compile the Model:

```

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

```

6. Train the Model:

```

history = model.fit(x_train, y_train, epochs=5, batch_size=64, validation_data=(x_test,
y_test))
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_acc}')

```

8. Plot Training and Validation Accuracy (Optional):

```

plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

```

Explanation of Code

- Conv2D Layers: These layers perform the convolution operation, where 32, 64, etc., represent the number of filters, and (3, 3) represents the size of each filter.
- MaxPooling2D: This reduces the spatial dimensions of the image (downsampling)

after each convolutional layer.

- Flatten: This layer flattens the 2D matrix output of the last convolutional layer into a 1D vector so that it can be passed to the dense layers.
- Dense Layers: These are fully connected layers that help make final predictions.
- Softmax Activation: This is used in the output layer for multi-class classification, where each output neuron corresponds to one class (digit 0-9 in this case).