

The IBM logo is displayed in its classic eight-striped font. The background of the top section of the slide is a dark blue field with faint, glowing white lines and shapes that suggest a complex network or data flow.

# Artificial Intelligence

Python

## Training Module



## Multi-Layer Perceptron (MLP) Overview

A Multi-Layer Perceptron (MLP) is a type of artificial neural network composed of multiple layers of neurons. It is a feedforward network, meaning the information flows in one direction from input to output, without any cycles or loops.

### Key components of an MLP:

1. Input layer: The first layer that takes the input features.
2. Hidden layers: Intermediate layers where the actual computation happens. An MLP typically has one or more hidden layers.
3. Output layer: The final layer that produces the prediction or output.

The neurons in each layer are connected to the neurons of the previous and next layers via weighted connections. Each neuron in a layer performs a weighted sum of the inputs, adds a bias term, and then passes this sum through an activation function, such as ReLU (Rectified Linear Unit) or sigmoid.

### Architecture

- Feedforward architecture: Information moves forward from the input layer to the output layer.
- Backpropagation: The process of training an MLP involves adjusting the weights using a method called backpropagation, which minimizes the error in the output by calculating the gradient of the error with respect to the weights.

### Steps in Training an MLP:

1. Forward propagation: Pass the input data through the network to obtain predictions.
2. Loss computation: Compute the loss (error) between the predicted output and the actual target.
3. Backward propagation: Compute the gradients of the loss with respect to each weight using backpropagation.
4. Weight update: Adjust the weights using an optimization algorithm like Gradient Descent or Adam.

### Python Code for a Basic MLP

Below is a simple implementation of an MLP for classification using Keras, a high-level neural network API in Python that runs on top of TensorFlow.

python

Copy code

```
# Import necessary libraries

import numpy as np

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score


# Load a dataset (Iris dataset for classification)
data = load_iris()
X = data.data
y = data.target


# Preprocess the data: scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)


# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)


# Define the MLP model
model = Sequential()

# Add input layer and first hidden layer with ReLU activation
model.add(Dense(64, input_dim=X_train.shape[1], activation='relu'))
```

```

# Add second hidden layer with ReLU activation
model.add(Dense(32, activation='relu'))

# Add output layer with softmax activation (for multi-class classification)
model.add(Dense(3, activation='softmax'))

# Compile the model with categorical crossentropy loss and Adam optimizer
model.compile(loss='sparse_categorical_crossentropy', optimizer=Adam(),
metrics=['accuracy'])

# Train the model
model.fit(X_train, y_train, epochs=50, batch_size=16, verbose=1)

# Evaluate the model on the test set
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1) # Convert probabilities to class labels

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred_classes)
print(f"Test accuracy: {accuracy:.2f}")

```

Explanation of the Code:

#### 1. Data Preparation:

- We use the Iris dataset from `sklearn.datasets`, which is a famous dataset for classification.
- The features are scaled using `StandardScaler` to ensure the network trains efficiently, as neural networks tend to perform better when the input data is standardized.
- The data is split into training and testing sets using `train_test_split`.

#### 2. Model Architecture:

- We create a `Sequential` model, which allows you to build a linear stack of layers.

- The first layer (Dense) is the input layer with 64 neurons and uses the ReLU activation function. The `input_dim` is set to the number of features in the input data (4 in this case for Iris).
- The second layer has 32 neurons and uses the ReLU activation function as well.
- The output layer consists of 3 neurons (one for each class in the Iris dataset) and uses the Softmax activation function to output probabilities for each class.

### 3. Compilation:

- The model is compiled with categorical cross-entropy loss (since this is a multi-class classification problem) and the Adam optimizer, which adapts the learning rate during training.

### 4. Training:

- The model is trained for 50 epochs, using a batch size of 16. This means the model updates its weights after processing 16 training samples.

### 5. Prediction and Evaluation:

- After training, the model predicts the classes for the test data. We use `argmax` to convert the output probabilities to predicted class labels.
- The accuracy of the model is calculated using the `accuracy_score` function.

### MLP Variations:

- Number of layers and neurons: You can adjust the number of hidden layers and neurons depending on the complexity of the task.
- Activation Functions: You can experiment with different activation functions (e.g., ReLU, sigmoid, tanh) based on your data and problem type.
- Optimization: The model can also be optimized with different algorithms, such as Stochastic Gradient Descent (SGD) or Adam.